

Evaluating performance of STL containers (unsorted)

Operation 1: End Insertion

In order of speed (from fastest to slowest):

- vector
- list
- unordered_set
- set

For end insertion, vector was the fastest and set was the slowest. The reason why vector was the fastest for this operation is because the capacity of a vector is initially double its size, so the memory (which is contiguous) for the new element has already been allocated. When adding to the end of a linked list, in addition to allocating memory for the element, the relevant pointers to the next and previous element must be reassigned/initialized. The reason why it's slowest when adding to a set is because a set is typically implemented as a binary search tree and stores its elements by following a specific order. In an unordered_set, the elements are not sorted in any particular order, but organized into buckets depending on their hash values. Therefore, insertion into an unordered_set is faster than insertion into a set which is ordered.

Operation 2: Beginning Insertion

In order of speed (from fastest to slowest):

- list
- unordered_set
- set
- vector

For beginning insertion, list was the fastest and vector was the worst by far. In order to add elements to the beginning of a vector, each of the existing elements must be shifted to the next position, which therefore causes the insertion process to be the slowest of each of the four containers. As mentioned for the previous insertion operation, insertion into an unordered_set is faster than insertion into a set which is ordered. Beginning insertion is the fastest for a list, because all that is required is memory allocation for the new element and the reassignment/initialization of the relevant previous and next pointers for the current head and the new head.

Operation 3: Find 10K Numbers



In order of speed (from fastest to slowest):

- `unordered_set`
- `set`
- `vector`
- `list`

For finding numbers, `unordered_set` was the fastest, and `list` was the slowest. An `unordered_set` is fastest because its elements are retrieved based on their hash values. A `set` is faster than a `vector` or `list` because its elements are stored as a binary search tree. For both a `vector` and a `list`, each of the elements must be searched until the correct one is found, and so that results in the slowest speed. Although time complexity is $O(n)$ for both a `list` and a `vector`, searching a `vector` is slightly more efficient than a `list` because its elements are stored in contiguous memory so the next record is already in the cache, and the pointer can simply be easily advanced, while in a `list` the next pointer must be fetched.

Key takeaways

- `Vector` is the slowest (by orders of magnitude) for beginning insertions.
- If your program requires mostly insertion operations (where position matters, i.e. elements can't simply be added to the end), and very few lookup operations, use a `list` (or `unordered_set`) and definitely avoid `vectors`.
- If your program requires multiple lookups, avoid `lists` and `vectors`, and use an `unordered_set`. A `set` is second best, with an average time complexity of $O(\log n)$ compared to the $O(1)$ time complexity of an `unordered_set`.
- For a `set`, time complexity for insertion is roughly equivalent for both beginning and end.
- For an `unordered_set`, time complexity for insertion is roughly equivalent for both beginning and end.
- A `list` and an `unordered_set` have a roughly equivalent time complexity for insertions, but an `unordered_set` is order of magnitudes faster when used for lookup.

Below are the execution times of these operations on my computer, as well as the average time complexities for these 4 containers.

Execution time (ms)	End Insertion	Beginning Insertion	Find 10K Numbers
<code>std::vector</code>	7,000	2,250,068	3,164,803
<code>std::list</code>	13,961	20,155	4,644,959
<code>std::set</code>	74,800	88,575	2,992
<code>std::unordered_set</code>	39,894	32,153	998

Key: **Best**
 Worst

Avg Time Complexity	End Insertion	Beginning Insertion	Find 10K Numbers
<code>std::vector</code>	$O(1)$	$O(n)$	$O(n)$
<code>std::list</code>	$O(1)$	$O(1)$	$O(n)$
<code>std::set</code>	$O(\log n)$	$O(\log n)$	$O(\log n)$
<code>std::unordered_set</code>	$O(1)$	$O(1)$	$O(1)$

Some sample outputs:

```
Execution time for inserting at end:
CONTAINER      TIME (ms)      PERCENTAGE
Vector-----7e+003-----100.00%
List-----13961-----199.96%
Set-----74800-----1071.33%
Unordered Set-----39894-----571.38%

Execution time for inserting at beginning:
CONTAINER      TIME (ms)      PERCENTAGE
Vector-----2250068-----100.00%
List-----20155-----0.90%
Set-----88575-----3.94%
Unordered Set-----32153-----1.43%

Execution time for finding 10,000 numbers:
CONTAINER      TIME (ms)      PERCENTAGE
Vector-----3164803-----100.00%
List-----4644959-----146.77%
Set-----2992-----0.09%
Unordered Set-----998-----0.03%
```

```
Execution time for inserting at end:
CONTAINER      TIME (ms)      PERCENTAGE
Vector-----9e+003-----100.00%
List-----14962-----166.69%
Set-----80783-----899.99%
Unordered Set-----39893-----444.44%

Execution time for inserting at beginning:
CONTAINER      TIME (ms)      PERCENTAGE
Vector-----2127425-----100.00%
List-----10970-----0.52%
Set-----78826-----3.71%
Unordered Set-----32876-----1.55%

Execution time for finding 10,000 numbers:
CONTAINER      TIME (ms)      PERCENTAGE
Vector-----2997822-----100.00%
List-----4238235-----141.38%
Set-----2013-----0.07%
Unordered Set-----0-----0.00%
```

```
Execution time for inserting at end:
CONTAINER      TIME (ms)      PERCENTAGE
Vector-----8e+003-----100.00%
List-----14779-----193.92%
Set-----104055-----1365.37%
Unordered Set-----41917-----550.02%

Execution time for inserting at beginning:
CONTAINER      TIME (ms)      PERCENTAGE
Vector-----2209445-----100.00%
List-----10662-----0.48%
Set-----85158-----3.85%
Unordered Set-----31620-----1.43%

Execution time for finding 10,000 numbers:
CONTAINER      TIME (ms)      PERCENTAGE
Vector-----2592631-----100.00%
List-----4823027-----186.03%
Set-----3535-----0.14%
Unordered Set-----1619-----0.06%
```